

Comprehensive Model Report

Comparative Analysis of Deep Learning Models for EuroSAT Classification

This report provides a detailed comparison of three distinct Convolutional Neural Network (CNN) models used for classifying images from the EuroSAT dataset: a fine-tuned VGG16, a custom CNN with the Swish activation function, and a fine-tuned ResNet50.

1. VGG16 with Fine-Tuning

Model Architecture and Preprocessing

- Base Model:** The VGG16 architecture, pre-trained on ImageNet, was used as the feature extraction base. The top classification layer of the original model was excluded (`include_top=False`).
- Custom Head:** A new classification head was added on top of the base model. It consists of a `Flatten` layer, a `Dense` layer with 256 neurons and ReLU activation, a `Dropout` layer with a rate of 0.5 for regularization, and a final `Dense` output layer with softmax activation for the number of classes in the dataset.
- Input Shape:** The model was configured to accept input images of size `64x64x3`.
- Preprocessing:** Input images were resized to `64x64` and their pixel values were normalized by casting to `tf.float32` and scaling to a `[0.0, 1.0]` range.

Training Strategy

- Two-Phase Training:**
 - Initial Training:* The model was first trained for 10 epochs with the entire VGG16 base frozen (`base_model.trainable = False`).
 - Fine-Tuning:* The VGG16 base was made trainable (`base_model.trainable = True`), but only the top 4 layers were unfrozen (`for layer in base_model.layers[:-4]: layer.trainable = False`).
- Optimizer:** Adam (`1e-4` for initial training, `1e-5` for fine-tuning)
- Loss Function:** `sparse_categorical_crossentropy`
- Batch Size:** 32
- Epochs:** Up to 30 (early stopping applied)
- Callbacks:**
 - `EarlyStopping(patience=5)`
 - `ReduceLROnPlateau(patience=3)`

Performance

- Initial Validation Accuracy:** 86.76%
- Fine-Tuned Validation Accuracy:** 93.04%

2. Custom CNN with Swish Activation

Model Architecture and Preprocessing

- Architecture:** A sequential model built from scratch, composed of three convolutional blocks followed by a dense classifier.
- Conv Blocks:** Each block includes two `Conv2D` layers with Swish activation, `BatchNormalization`, `MaxPooling2D`, and `Dropout` (rate 0.3). Filter sizes increased across blocks: `32 → 64 → 128`.
- Classifier Head:** `Flatten → Dense(256, Swish) → BatchNormalization → Dropout(0.5) → Dense softmax output`
- Input Shape:** `128x128x3`
- Data Augmentation:** `RandomFlip`, `RandomRotation`, and `RandomZoom`
- Preprocessing:** Images resized to `128x128`, scaled to `[0.0, 1.0]`

Training Strategy

- Optimizer:** `RMSprop(learning_rate = 0.001)`
- Loss Function:** `sparse_categorical_crossentropy`
- Batch Size:** 32
- Epochs:** Up to 50
- Callbacks:**
 - `EarlyStopping(patience=7, monitor='val_accuracy')`
 - `ReduceLROnPlateau(patience=3)`
 - `ModelCheckpoint`

Performance

- Peak Validation Accuracy:** 95.56%

3. ResNet50 with Fine-Tuning

Model Architecture and Preprocessing

- **Base Model:** Pre-trained ResNet50 (`include_top=False`)
- **Custom Head:** `GlobalAveragePooling2D` → `Dense(256, ReLU)` → `Dropout(0.5)` → softmax output
- **Input Shape:** 224x224x3
- **Preprocessing:** Used `tf.keras.applications.resnet.preprocess_input` ; labels were one-hot encoded

Training Strategy

- **Two-Phase Training:**
 - *Initial Training:* 15 epochs with frozen base
 - *Fine-Tuning:* Top 10 layers unfrozen (`for layer in base_model.layers[:-10]: layer.trainable = False`), trained for 10 more epochs
- **Optimizer:** Adam (`1e-4` then `1e-5`)
- **Loss Function:** `categorical_crossentropy`
- **Batch Size:** 32
- **Callback:** `ModelCheckpoint(monitor='val_accuracy')`

Performance

- **Initial Validation Accuracy:** 96.96%
- **Fine-Tuned Validation Accuracy:** 97.04%

Model Comparison Table

Feature	VGG16 with Fine-Tuning	Custom CNN (with Swish)	ResNet50 with Fine-Tuning
Base Architecture	Pre-trained VGG16	Custom Sequential Model	Pre-trained ResNet50
Input Shape	64x64x3	128x128x3	224x224x3
Data Preprocessing	Resize, scale to [0, 1]	Resize, scale to [0, 1] , data augmentation	Resize, <code>resnet.preprocess_input</code>
Optimizer	Adam (LR: <code>1e-4</code> → <code>1e-5</code>)	RMSprop (LR: <code>0.001</code>)	Adam (LR: <code>1e-4</code> → <code>1e-5</code>)
Loss Function	<code>sparse_categorical_crossentropy</code>	<code>sparse_categorical_crossentropy</code>	<code>categorical_crossentropy</code>
Training Strategy	Fine-tuned top 4 layers after initial training	Trained from scratch	Fine-tuned top 10 layers after initial training
Peak Validation Accuracy	93.04%	95.56%	97.04%